

# Assignment 1

## Q1

F1 instances include 16 nm Xilinx UltraScale Plus FPGA.

### F1 FPGA specification

| <u>Aa</u> Name                | <u>≡</u> Amount |
|-------------------------------|-----------------|
| <u>SLR</u>                    | 3               |
| <u>System Logic Cells (K)</u> | 2,586           |
| <u>CLB LUTs (K)</u>           | 1,182           |
| <u>CLB LUTs (K)</u>           | 1,182           |
| <u>DSP Slices</u>             | 6,840           |
| <u>URAM</u>                   | 270.0Mb         |
| <u>Total Block RAM (Mb)</u>   | 75.9            |
| <u>DDR total capacity</u>     | 64GB (4x16GB)   |
| <u>DDR Total BW</u>           | 68GB/s          |
| <u>PCI Express</u>            | Gen3x16         |

### F1 CPU specification

```
Architecture:      x86_64
CPU op-mode(s):    32-bit, 64-bit
Byte Order:        Little Endian
CPU(s):            8
On-line CPU(s) list: 0-7
Thread(s) per core: 2
Core(s) per socket: 4
Socket(s):         1
NUMA node(s):      1
Vendor ID:         GenuineIntel
CPU family:        6
Model:             79
Model name:        Intel(R) Xeon(R) CPU E5-2686 v4 @ 2.30GHz
Stepping:          1
CPU MHz:           1497.788
CPU max MHz:       3000.0000
CPU min MHz:       1200.0000
BogoMIPS:          4600.00
Hypervisor vendor: Xen
Virtualization type: full
L1d cache:         32K
L1i cache:         32K
L2 cache:          256K
L3 cache:          46080K
NUMA node0 CPU(s): 0-7
```

## Q2

The PCIe report state that

```
[ec2-user@ip-172-31-17-89 vitis_intro_lab]$ xbutil dmatest
-----
Deprecation Warning:
  The given legacy sub-command and/or option has been deprecated
  to be obsoleted in the next release.

  Further information regarding the legacy deprecated sub-commands
  and options along with their mappings to the next generation
  sub-commands and options can be found on the Xilinx Runtime (XRT)
  documentation page:

  https://xilinx.github.io/XRT/master/html/xbtools_map.html

  Please update your scripts and tools to use the next generation
  sub-commands and options.
-----
INFO: Found total 1 card(s), 1 are usable
INFO: DMA test on [0]: xilinx_aws-vu9p-f1_shell-v04261818_201920_2
Total DDR size: 65536 MB
Buffer Size: 16 MB
Reporting from mem_topology:
Data Validity & DMA Test on bank0
Host -> PCIe -> FPGA write bandwidth = 8500.711852 MB/s
Host <- PCIe <- FPGA read bandwidth = 12200.716077 MB/s
INFO: xbutil dmatest succeeded.
```


## Q3

### A

The results are obtained using Chrono high-resolution clock. This result are obtained using the vadd code

#### Vadd CPU time vs FPGA time


| <u>Aa</u> vector size | <u>#</u> Data transfer time millisecond | <u>#</u> FPGA time (microsecond) | <u>#</u> CPU time microsecond |
|-----------------------|---|----------------------------------|-------------------------------|
| <u>1024</u>           | 22.759                                  | 150                              | 5                             |
| <u>2048</u>           | 21.146                                  | 192                              | 6                             |
| <u>4096</u>           | 21.972                                  | 214                              | 7                             |
| <u>8192</u>           | 22.893                                  | 287                              | 11                            |
| <u>16384</u>          | 21.955                                  | 344                              | 15                            |
| <u>32768</u>          | 22.334                                  | 550                              | 30                            |
| <u>65536</u>          | 22.545                                  | 983                              | 58                            |
| <u>131072</u>         | 23.856                                  | 1834                             | 114                           |
| <u>262144</u>         | 25.556                                  | 3525                             | 230                           |
| <u>524288</u>         | 26.462                                  | 6728                             | 492                           |
| <u>1048576</u>        | 27.561                                  | 14484                            | 960                           |
| <u>2097152</u>        | 28.122                                  | 25934                            | 2245                          |
| <u>4194304</u>        | 30.453                                  | 54112                            | 5277                          |
| <u>8388608</u>        | 33.315                                  | 107125                           | 10291                         |
| <u>16777216</u>       | 84.634                                  | 211362                           | 21576                         |
| <u>33554432</u>       | 103.225                                 | 422904                           | 48914                         |



In wide add we use uint512\_t which enable us to directly take 512 bits at a time. Also wide add use memory banking which help in reducing the congestion of reading the memory and different variables can read data in parallel.

#### wide add CPU time vs FPGA time

| vector size | Data transfer time millisecond | FPGA time (microsecond) | CPU time microsecond |
|-------------|--------------------------------|-------------------------|----------------------|
| 1024        | 21.449                         | 97                      | 5                    |
| 2048        | 22.156                         | 141                     | 6                    |
| 4096        | 21.562                         | 129                     | 7                    |
| 8192        | 22.693                         | 112                     | 11                   |
| 16384       | 22.065                         | 146                     | 15                   |
| 32768       | 22.235                         | 169                     | 30                   |
| 65536       | 21.652                         | 215                     | 58                   |
| 131072      | 23.452                         | 266                     | 114                  |
| 262144      | 24.856                         | 545                     | 230                  |
| 524288      | 26.462                         | 892                     | 492                  |
| 1048576     | 27.451                         | 1687                    | 960                  |
| 2097152     | 28.562                         | 3381                    | 2245                 |
| 4194304     | 29.953                         | 6464                    | 5277                 |
| 8388608     | 32.315                         | 12985                   | 10291                |
| 16777216    | 86.654                         | 25797                   | 21576                |
| 33554432    | 102.255                        | 51437                   | 48914                |




Here we can see that we are able to obtain better results compared to simple vadd code.

## B

Using the wide add code we modified it so that is able to read floating point data type and instead of addition we do multiplication.


float mult CPU time vs FPGA time

| <u>Aa</u> vector size | # Data transfer time millisecond | # FPGA time (microsecond) | # CPU time microsecond |
|-----------------------|----------------------------------|---------------------------|------------------------|
| <u>1024</u>           | 19.65                            | 105                       | 2                      |
| <u>2048</u>           | 19.145                           | 129                       | 2                      |
| <u>4096</u>           | 20.617                           | 133                       | 4                      |
| <u>8192</u>           | 19.627                           | 120                       | 7                      |
| <u>16384</u>          | 21.363                           | 145                       | 14                     |
| <u>32768</u>          | 20.988                           | 154                       | 47                     |
| <u>65536</u>          | 23.206                           | 217                       | 58                     |
| <u>131072</u>         | 23.548                           | 301                       | 114                    |
| <u>262144</u>         | 24.045                           | 447                       | 225                    |
| <u>524288</u>         | 26.696                           | 817                       | 455                    |
| <u>1048576</u>        | 30.811                           | 1541                      | 1011                   |
| <u>2097152</u>        | 32.68                            | 2951                      | 2244                   |
| <u>4194304</u>        | 38.241                           | 5820                      | 4816                   |
| <u>8388608</u>        | 41.911                           | 11541                     | 10855                  |
| <u>16777216</u>       | 62.401                           | 22981                     | 21400                  |
| <u>33554432</u>       | 75.247                           | 45926                     | 39931                  |



## Q4

Using the code of wide add instead of sending the data and then reading the result all at once we divide the data into n chunks. this pre




This help us in reducing the total time taken to compute the data

### Parallel CPU vs FPGA

| Array size | FPGA time (ms) | Communication time (ms) | CPU time(ms) |
|------------|----------------|-------------------------|--------------|
| 32768      | 0.74           | 20.34                   | 0.105        |
| 65536      | 0.761          | 19.817                  | 0.263        |
| 131072     | 0.752          | 20.976                  | 0.521        |
| 262144     | 0.809          | 22.369                  | 0.832        |
| 524288     | 0.743          | 23.57                   | 0.812        |
| 1048576    | 1.615          | 28.952                  | 1.418        |
| 2097152    | 3.368          | 30.03                   | 3.857        |
| 4194304    | 6.026          | 35.132                  | 4.886        |
| 8388608    | 11.461         | 43.015                  | 9.488        |
| 16777216   | 22.264         | 75.487                  | 18.929       |
| 33554432   | 44.067         | 135.372                 | 36.941       |
| 67108864   | 87.6           | 254.892                 | 74.462       |

| Aa Array size    | # FPGA time (ms) | # Communication time (ms) | # CPU time(ms) |
|------------------|------------------|---------------------------|----------------|
| <u>134217724</u> | 176.932          | 450.909                   | 146.944        |



## Q5

Kernel code

```
#include <stdio.h>
#include <string.h>

extern "C"
{
    void matrix_multiply(float *data, float *query, int N, float *res)
    {
        #pragma HLS INTERFACE m_axi port = data max_read_burst_length = 32 offset = slave bundle = gmem
        #pragma HLS INTERFACE m_axi port = query max_read_burst_length = 32 offset = slave bundle = gmem1
        #pragma HLS INTERFACE m_axi port = res max_write_burst_length = 32 offset = slave bundle = gmem2
        #pragma HLS INTERFACE s_axilite port = data bundle = control
        #pragma HLS INTERFACE s_axilite port = query bundle = control
        #pragma HLS INTERFACE s_axilite port = res bundle = control
        #pragma HLS INTERFACE s_axilite port = N bundle = control
        #pragma HLS INTERFACE s_axilite port = return bundle = control
        float query_local[256]; // 4 queries
        float query_cal;
        // float data_val[256];
        int n = N;
        Load_Query: for (int i = 0; i < 256; i++)
        {
            #pragma HLS pipeline
            #pragma HLS unroll factor=16
            query_local[i] = query[i] ;
        }

        Main:
        for (int j = 0; j < n; j++)
        {
            // #pragma HLS pipeline
            #pragma HLS LOOP_TRIPCOUNT min = 16 max = 1024

            query_cal = 0 ;

            M_cp:
            for (int k = 0; k < 256; k++)
            {
                #pragma HLS pipeline
                #pragma HLS unroll factor=16

                query_cal += data[k + j * 256] * query_local[k] ;
            }
        }
    }
}
```

```

        res[j] =query_cal ;
    }
}
}

```

## Host code

```

#include <vector>
#include <random>
#include "xcl2.hpp"
#include <algorithm>
// #include <cstring>
#include <string>
#include <iostream>
#include <chrono>

using namespace std;

int getans(float *data , float *query , int N ){
    int ans = 0;
    float val = 0 ;
    std::cout<< N << std::endl ;
    for (int i =0 ; i< N ;i++){
        float tem = 0 ;
        for (int j= 0 ; j < 256 ; j++){
            tem += data[i* 256 + j ] * query[j] ;
        }
        if (tem > val){
            val = tem;
            ans = i ;
        }
        //std::cout << tem << std::endl;
    }
    return ans ;
}

int subdivide_buffer(std::vector<cl::Buffer> &data_bufs, cl::Buffer data_in,
    std::vector<cl::Buffer>&res_bufs , cl::Buffer res_in , uint &num_div , int N){
    num_div = (N +511)/512;
    cl_buffer_region r1, r2 ;
    int err;
    r1.origin = 0 ;
    r2.origin = 0 ;

    r1.size = 256 * 4*16;
    r2.size = 4 *16 ;
    for(uint i =0 ;i< num_div ; i++){
        if (i == num_div -1 ){
            r1.size =data_in.getInfo<CL_MEM_SIZE>() - i*16*4*256;
            r2.size =res_in.getInfo<CL_MEM_SIZE>() - i*16*4;
        }
        data_bufs.push_back(data_in.createSubBuffer(static_cast<cl_mem_flags>(CL_MEM_READ_ONLY) , CL_BUFFER_CREATE_TYPE_REGION, &r1, &err))
        if (err != CL_SUCCESS){
            exit(-1) ;
        }
        res_bufs.push_back(res_in.createSubBuffer(static_cast<cl_mem_flags>(CL_MEM_WRITE_ONLY), CL_BUFFER_CREATE_TYPE_REGION, &r2, &err));
        if (err != CL_SUCCESS){
            exit(-1) ;
        }
        r1.origin += r1.size;
        r2.origin += r2.size ;
    }
    return 0 ;
}

int enqueue_subbuf_vadd(cl::CommandQueue &q, cl::Kernel &krnl,
    cl::Event &event, cl::Buffer a, cl::Buffer b, cl::Buffer c)
{
    // Get the size of the buffer
    cl::Event k_event, m_event;
    std::vector<cl::Event> krnl_events;

    static std::vector<cl::Event> tx_events, rx_events;
    std::vector<cl::Memory> c_vec;
    size_t size;
    size = a.getInfo<CL_MEM_SIZE>();

    std::vector<cl::Memory> in_vec;
    in_vec.push_back(a);
    in_vec.push_back(b);
    q.enqueueMigrateMemObjects(in_vec, 0, &tx_events, &m_event);
    krnl_events.push_back(m_event);
    tx_events.push_back(m_event);

```

```

    if (tx_events.size() > 1)
    {
        tx_events[0] = tx_events[1];
        tx_events.pop_back();
    }

    krnl.setArg(0, a);
    krnl.setArg(1, b);
    krnl.setArg(3, c);
    krnl.setArg(2, (int)(size / (4 * 256)));
    q.enqueueTask(krnl, &krnl_events, &k_event);
    krnl_events.push_back(k_event);
    if (rx_events.size() == 1)
    {
        krnl_events.push_back(rx_events[0]);
        rx_events.pop_back();
    }
    c_vec.push_back(c);
    q.enqueueMigrateMemObjects(c_vec, CL_MIGRATE_MEM_OBJECT_HOST, &krnl_events, &event);
    rx_events.push_back(event);

    return 0;
}

int main(int argc, char **argv)
{
    if (argc < 2 || argc > 4)
    {
        cout << "Incorrect format" << endl;
        return EXIT_FAILURE;
    }
    char *binaryFile = argv[1];
    int N ;
    if (argc == 3)
    {
        try
        {
            N = stoi(argv[2]);
        }
        catch (invalid_argument val)
        {
            cerr << "Invalid argument" << endl;
            return -1;
        }
    }
    else
    {
        N = 256;
    }

    cl::Device device;
    cl::Context context;
    cl::CommandQueue q;
    cl::Program program;
    cl::Kernel krnl;
    cl_int err;

    auto devices = xcl::get_xil_devices();

    auto fileBuf = xcl::read_binary_file(binaryFile);
    cl::Program::Binaries bins{{fileBuf.data(), fileBuf.size()}};
    // bool valid_device = false;
    for (unsigned int i = 0; i < devices.size(); i++)
    {
        device = devices[i];
        // Creating Context and Command Queue for selected Device
        OCL_CHECK(err, context = cl::Context(device, NULL, NULL, NULL, &err));
        OCL_CHECK(err,
            q = cl::CommandQueue(context, device,
                                CL_QUEUE_PROFILING_ENABLE |
                                CL_QUEUE_OUT_OF_ORDER_EXEC_MODE_ENABLE,
                                &err));
        std::cout << "Trying to program device[" << i
                    << "]: " << device.getInfo<CL_DEVICE_NAME>() << std::endl;
        cl::Program program(context, {device}, bins, NULL, &err);
        if (err != CL_SUCCESS)
        {
            std::cout << "Failed to program device[" << i << "] with xclbin file!\n";
        }
        else
        {
            std::cout << "Device[" << i << "]: program successful!\n";
            // Creating Kernel
            OCL_CHECK(err, krnl = cl::Kernel(program, "matrix_multiply", &err));
            // valid_device = true;
            break; // we break because we found a valid device
        }
    }
}

```



```

    }
}

cl::Buffer data_buf(context,
    static_cast<cl_mem_flags>(CL_MEM_READ_ONLY),
    N*256 * sizeof(float),
    NULL,
    NULL);
cl::Buffer query_buf(context,
    static_cast<cl_mem_flags>(CL_MEM_READ_ONLY),
    256 * sizeof(float),
    NULL,
    NULL);
cl::Buffer res_buf(context,
    static_cast<cl_mem_flags>(CL_MEM_READ_WRITE),
    N* sizeof(float),
    NULL,
    NULL);
krnl.setArg(0, data_buf);
krnl.setArg(1, query_buf);
krnl.setArg(2, N);
krnl.setArg(3, res_buf);
float *data = (float *)q.enqueueMapBuffer(data_buf,
    CL_TRUE,
    CL_MAP_WRITE,
    0,
    N*256 * sizeof(float));
float *query = (float *)q.enqueueMapBuffer(query_buf,
    CL_TRUE,
    CL_MAP_WRITE,
    0,
    256* sizeof(float));

// Filling the data
float temp_sum = 0;
for (int i = 0; i < N; i++)
{
    temp_sum = 0 ;
    for (int j = 0; j < 256; j++)
    {
        data[i * 256 + j] = static_cast<float>(rand()) / static_cast<float>(RAND_MAX);
        temp_sum += data[i * 256 + j] * data[i * 256 + j];
    }
    temp_sum = sqrt(temp_sum);
    for (int j = 0; j < 256; j++)
    {
        data[i * 256 + j] /= temp_sum;
    }
}
temp_sum = 0 ;
for (int j = 0; j < 256; j++)
{
    query[j] = static_cast<float>(rand()) / static_cast<float>(RAND_MAX);
    temp_sum += query[j]*query[j];
}
temp_sum = sqrt(temp_sum);
for (int j = 0; j < 256; j++)
{
    query[j] /= temp_sum;
}

std::chrono::high_resolution_clock::time_point st = std::chrono::high_resolution_clock::now() ;
int ans_val = getans(data , query , N);
std::chrono::duration<float, std::milli> duration = std::chrono::high_resolution_clock::now() - st;

std::cout << "CPU time:\t" << duration.count() << " ms " << endl;
q.enqueueUnmapMemObject(data_buf, data);
q.enqueueUnmapMemObject(query_buf, query);

std::vector<cl::Buffer> data_bufs , res_bufs;
uint num_div;
subdivide_buffer(data_bufs, data_buf, res_bufs, res_buf, num_div, N);
cl::Event kernel_events[num_div];

for (uint i = 0; i < num_div; i++)
{
    cl::Buffer t2 = data_bufs[i];
    cl::Buffer te = res_bufs[i] ;
    enqueue_subbuf_vadd(q, krnl, kernel_events[i], t2 , query_buf, te);
}
st = std::chrono::high_resolution_clock::now() ;
clWaitForEvents(num_div, (const cl_event *)&kernel_events);
duration = std::chrono::high_resolution_clock::now() - st;

```

```

std::cout << "FPGA time:\t" << duration.count() << " ms " << endl;
int ker_ans = 0, tem_ans = 0;
float *res = (float *)q.enqueueMapBuffer(res_buf,
                                         CL_TRUE,
                                         CL_MAP_READ,
                                         0,
                                         N * sizeof(float));

//calculate ans
for (int i = 0; i < N; i++){
    if (tem_ans < res[i]){
        tem_ans = res[i];
        ker_ans = i;
    }
}

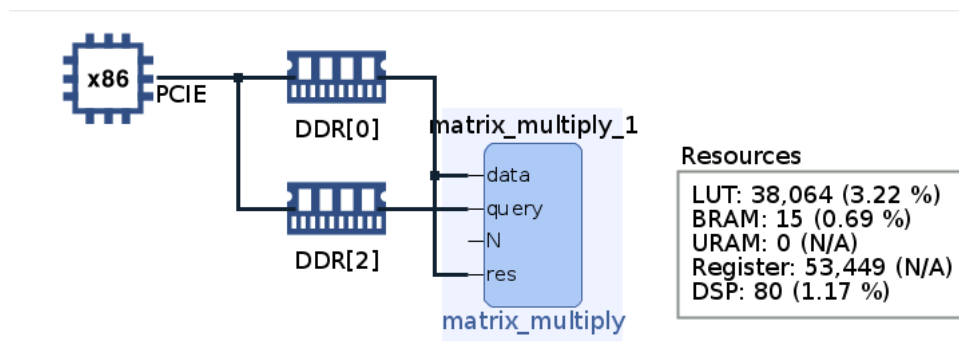
q.enqueueUnmapMemObject(res_buf, res);

std::cout << "Finish \n";
return 0;
}

```


Let us go through the code

- We see that computing dot product is the most compute-intensive task and we create a kernel for it.
- First, we load the query on the local memory for fast io. (Load\_Query)
- As we know that at max 512 bits can be transferred at a time we unroll the loop by a factor of 16 and then pipeline it.
- To reduce the data transfer time we divide the buffer into 512 size chunks.
- We also bank the data in different ddr so that we can access them in parallel.



#### Code stastics

| <a href="#">Aa</a> Database size | ⌘ Latency (ms) | ⌘ Throughput (Mb/s) |
|----------------------------------|----------------|---------------------|
| <a href="#">2048</a>             | 0.864931       | 9.03251             |
| <a href="#">4096</a>             | 1.85606        | 8.41835             |
| <a href="#">8192</a>             | 3.58039        | 8.7281              |
| <a href="#">16384</a>            | 7.16698        | 8.72055             |
| <a href="#">32768</a>            | 13.4871        | 9.26808             |
| <a href="#">65536</a>            | 28.9076        | 8.64824             |
| <a href="#">131072</a>           | 54.8245        | 9.12001             |
| <a href="#">262144</a>           | 109.119        | 9.16434             |
| <a href="#">524288</a>           | 215.357        | 9.28689             |
| <a href="#">1048576</a>          | 425.91         | 9.39165             |



## Roofline analysis

For now, let us assume all floating-point multiplication and accumulation happen in DSPs. In the AWS F1 instance, the FPGA has 6800 DSPs. These DSPs compute 1 Mac operating in 1 clock cycle.

Clock frequency = 250 MHz

Peak computation performance =  $6800 * 250 * 10^6 = 1.7$  TFlops

The DDR memory bandwidth = 16 GB/s

The Operational Intensity =  $2(\text{multiplication} + \text{addition}) / 4$  (4 as we are reading or writing 4byte in parallel) =  $1/2$

Attainable GFlops =  $1/2 * 16 = 8$  GFlops

So the code is memory-bound as 8GFlops is very less compared to 1.7 TFlops this is assuming we are doing the best designing we can.

The current avg throughput is 8.977 Mb/s